

# Table of Contents

<b>Chapter 1: Background</b>	<b>3</b>
1.1 Analyzing Algorithms	3
1.1.1 Big-O	3
1.1.2 P vs NP	5
1.2 Graph Theory	6
1.3 Mechanism Design	8
1.3.1 Stable Matching	8
1.3.2 Fair Division	10
1.4 Adding Constraints	11
1.4.1 Quotas and Diversity	11
1.4.2 Optimization	12
1.5 Related Work	12
1.5.1 Two-sided matching	13
1.5.2 One-sided matching	13



# Chapter 1

## Background

### 1.1 Analyzing Algorithms

An algorithm is a list of specific instructions written to achieve some goal. The goal varies, and in casual conversation individuals have come to associate *algorithm* with the phrase “my algorithm” or “the algorithm,” an entity with the goal to create profit for big tech companies by keeping users engaged on their apps. While those algorithms exist, they are only one type of algorithm. A cookie recipe is an algorithm, providing instructions a baker follows to yield a batch of cookies upon completion. GPS devices provide algorithms for driving from one location to another. In computer science, an algorithm is a sequence of lines of code that perform a computation, usually mathematical calculations or processing data. Because algorithms are so versatile, and there are so many different ways of accomplishing one goal, we analyze algorithms by their efficiency: how many resources they use and how much time they take to complete the task.

#### 1.1.1 Big-O

To analyze efficiency, we have to define what we mean for an algorithm to be “efficient.” For the sake of this thesis, we will focus on the time aspect of efficiency. The amount of time an algorithm takes to execute from start to finish is known as the *running time* of an algorithm. We can mathematically formalize running time with what’s known as *Big-O* notation.

Big-O notation is a form of *asymptotic* analysis, or analysis that takes place as the size of the input becomes infinitely large. It is important to analyze algorithms asymptotically so that we know the algorithm is still effective in the worst case scenario. This allows us to guarantee an algorithm scales well for arbitrarily large inputs. For instance, if I am following an algorithm to bake cookies and the steps involve rolling out the dough, mixing, and measuring ingredients by hand, the algorithm works well for me to make a batch of 24 cookies. However, if I start a popular cookie franchise and suddenly I need 2,400 cookies a day, my original algorithm is not going to be the

most efficient way to make that many cookies. While it may make sense when baking cookies to handle these two cases differently, in many computer science problems data commonly scales to very large sizes, so it is helpful to know our algorithms will still be efficient in those cases.

We *estimate* running time through asymptotic analysis. In Big-O notation, we write that for some algorithm called  $f$ ,  $f(n) = O(g(n))$ . This is read “ $f(n)$  is big  $O$  of  $g(n)$ ,” where  $f$  on input  $n$  will take **at most**  $g(n)$  time to run. The big  $O$  signifies the “at most” in that statement, or more formally it denotes an *upper bound* on the asymptotic running time. Running time can be classified into different categories. The most relevant categories to this thesis are constant time, polynomial time, and exponential time. Constant time,  $O(1)$ , can be abstractly thought of as any operation that takes the same fixed amount of time regardless of the input size. Constant time is the fastest running time an algorithm can have. Most algorithms involve more complicated operations than those that can be done in constant time, however, so our next category of time analysis is polynomial time.

Recall that a polynomial is a mathematical expression that combines terms of variables and constants using various operations (e.g.  $x^4 + 2x^3 + 7$  or  $x^2 + y^2$ ). In time analysis, we simplify to the largest term in a polynomial and we disregard coefficients due to the fact that they become irrelevant as the input scales. For instance, say we have a number  $x^{100}$  versus  $3x^{100}$ . If  $x = 1,000$ , then both of these numbers are extremely large regardless of if we multiply it by 3. So in big-O notation, we would say  $f(x) = 3x^{100} = x^{100}$ . Similarly for multiple terms, in  $x^{100} + x^5$  the larger  $x^{100}$  term will be the part that grows bigger, so when analyzing time, it overshadows the  $x^5$  term. When we think about polynomial time, we think about how many actions must be completed per input to the algorithm. For instance, if I already have cookie dough and I need to bake the cookies, my algorithm might be: for every tablespoon of cookie dough, roll the dough into a ball and then sprinkle cinnamon on top. That is 2 actions to complete for every tablespoon of cookie dough. So, if we have  $n$  tablespoons, we could say this step takes  $O(2n) = O(n)$  time (note this would be more specifically linear time, but linear time is a subset of polynomial time). For our purposes, we can consider polynomial time to be sufficiently “efficient,” with the consideration that we would hope for our algorithm to be much closer to  $O(n)$  time than  $O(n^{100})$ , as the former is much more efficient than the latter. Our last category is exponential time, which is very slow compared to our other two categories. We consider an algorithm “inefficient” if it takes exponential time. An example of exponential time in Big-O is  $O(2^n)$ . As  $n$  grows larger in this case, the running time increases extremely rapidly. To compare polynomial time and exponential time, imagine we have some number  $n = 1000$ . If we square  $n$ , we get a polynomial expression:  $n^2 = 1,000,000$ . That is indeed large, but not compared to  $2^n$ , which is “a number much larger than the number of atoms in the universe” [1].

### 1.1.2 P vs NP

In time complexity theory, the field of mathematical analysis centered on the running time of algorithms, we give names to the aforementioned categories. Specifically, the most important complexity category for our purposes is P, which denotes the class of problems that can be solved with polynomial time algorithms. That means we can think of P as a collection of problems where every problem in the collection has an algorithm that will find a solution in polynomial time. Perhaps intuitively, we know there exist many complex problems that take longer than polynomial time to solve (such as those that take exponential time and belong to the class E). A less intuitive fact is that there also exist problems that we think do not belong in P, but we don't know how to prove it at this time nor do we believe we will know how to prove it anytime soon.

To understand what I mean by this, we need to discuss another important complexity class called NP: the class of problems solved by *nondeterministic polynomial* time algorithms. In very simplified terms, one could think of “nondeterministic” as a brute force approach: we take all possible solutions to a problem, and we have an algorithm that takes polynomial time to determine whether each solution actually solves the problem by going through them sequentially until it finds a solution that works. If the first solution happens to solve the problem, then the algorithm has been carried out once, so a polynomial number of operations have been performed. Therefore, the algorithm terminated in polynomial time. As an example, lets think about the classic 52-card game of solitaire. It is hard to know whether a starting configuration of cards will lead to a win in solitaire, but we can figure out if it is a winnable configuration by go through every single subsequent set of moves that leads to a terminating position and see if that final position is a win. If the first set of moves happens to lead to a win, we have moved a polynomial number of cards, and therefore we determined that the configuration is winnable in polynomial time [2].

It is not possible to guarantee that the first evaluated solution to an algorithm will always solve the problem, but we can imagine that if there were  $n$  possible solutions to an algorithm and we happened to have  $n$  computers that can all run the algorithm on a distinct solution, one of those computers would run the lucky input and solve the problem in polynomial time (if a valid solution exists). We could not determine which computer it would be, but we would know it exists. In real life, we do not have  $n$  computers for every problem we wish to solve with  $n$  possible solutions. However, if we were somehow given the solution output by that lucky computer, we could *verify* that it is correct by running the lucky computer's input on the real computer we do have in polynomial time. All of this is to say that NP is the class of problems that have algorithms for which we can verify a given solution is correct in polynomial time. We cannot necessarily find that solution in polynomial time, but if it was given to us we can prove it solves the problem in polynomial time.

Keep in mind that problems we know are in P are also in NP, because we can verify that their solutions (which we can find in polynomial time) are correct in polynomial

time. We will denote the problems that we cannot prove are or are not in  $P$  denote as  $NP \setminus P$ , which means the set difference of  $NP$  and  $P$ , or all problems in  $NP$  that are not known to be in  $P$ . Generally, we consider the class of problems in  $NP \setminus P$  to be very difficult to solve, as it means no one has ever discovered an algorithm that can find a correct solution in polynomial time. The hardest NP problems are known as *NP-complete* problems, where “complete” means the problem can be generalized to take the form of every other problem in NP. So, if a polynomial time algorithm was discovered for an NP-complete problem, it would mean that every NP problem could be solved in polynomial time, and therefore all of those problems would actually just be in  $P$  [1].

## 1.2 Graph Theory

In discrete mathematics, a graph is a collection of nodes called *vertices* connected by *edges* (denoted  $G = (V, E)$  where  $G$  is the graph, and it is composed of a set of vertices  $V$  and a set of edges  $E$ ). Mathematicians like graphs because their structure makes them very effective for modeling a variety of real-world situations such as social networks (where vertices are individuals, and edges connect people to their acquaintances) or even the spread of disease (where vertices are people and edges represent an infection from one person to another).

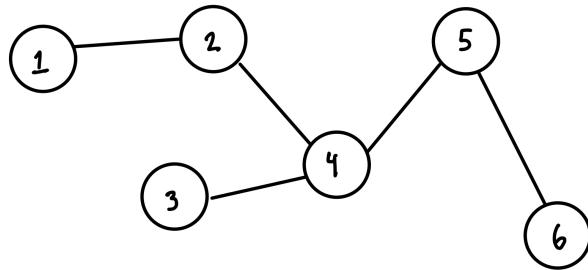


Figure 1.1: Simple graph with 6 vertices and 5 edges

Graphs can famously be used to model scheduling problems by adjusting their structure slightly. A basic graph problem has vertices and edges, but we can add factors such as *direction* going from one vertex to another, or some number of *colors* with which we want to color vertices such that no two adjacent vertices (vertices connected by an edge) are the same color. The latter example is actually one of the most studied NP-complete problems: is it possible to color a graph  $G$  using  $k$  colors? It is also the structure of graph under which scheduling has been extensively studied, where vertices represent events, edges encode overlapping time conflicts between events, and the number of colors is the number of available time slots into which those events can be scheduled [3].




Figure 1.2: Graph from Figure 1.1 colored with 3 colors

Another way to construct graphs is to add *weights* to edges. A weight in this context is some value derived based on the constraints of the problem, and this form of graph often lends itself to optimization. As in, the weighted graph can be used to model problems where we want to maximize or minimize the weights. For instance, we can construct a graph where each vertex represents a city and each edge represents a highway from one city to another. If we were planning a roadtrip, we could add a weight to each edge to signify the travel time of that highway from one city to another. Then, if we were planning a road trip to all the cities on the graph, we could find the fastest route by finding the path with the smallest weight.




Figure 1.3: Graph from Figure 1.1 with weights

## 1.3 Mechanism Design

Mechanism design has traditionally been studied by economists and game theorists, though in recent years computer scientists have increasingly joined the field because mechanisms are often just algorithms. Mechanism design is the construction of systems that “fairly” (under various definitions of the word) address scenarios involving a host of participants with adverse or overlapping goals. Some examples of problems that fall under this classification are as follows: designing a voting system that successfully eliminates strategizing by any party [4], auctioning search engine advertisement spots in a way that promotes honest bidding [5], assigning roommates from a pool of students such that each student ends up with a roommate they prefer [6], and maximizing aid while allocating food donations to food banks [7].

One of the most standard metrics for evaluating the effectiveness of a mechanism is whether it is *strategy-proof*, or prevents participants from taking advantage of the rules of the mechanism for their own gain. Many more metrics exist for evaluating subcategories of mechanism design, but almost all mechanisms can be evaluated under strategy-proofness. There are two subcategories of mechanism design relevant to this thesis: stable matching and fair division.

### 1.3.1 Stable Matching

Many real world scenarios can be modelled as matching problems, such as assigning medical students to hospitals for residency [8], or the aforementioned roommates example (known as the Stable Roommates problem). The standard (and heteronormative) matching problem is known as the Stable Marriage Problem [9], in which the goal is to match an equal number of men and women such that no participant remains uncoupled. This is known as a *bipartite* matching problem, or one in which the market of participants is divided into two disjoint sides and matched from one side to the other. Each participant provides a list of preferences over the other side, where they rank those whom they would most like to be matched to. These are known as *two-sided* preferences. An essential aspect of the Stable Marriage Problem is that the matching should prevent “cheating,” which occurs when a man and a woman who weren’t matched to each other prefer to be together over their assigned partners and will “elope,” leaving the other two participants alone. This couple would be known as a *blocking pair*, because their elopement “blocks” the matching from being complete. Herein lies the stability aspect of stable matching: we wish to find an algorithm that outputs a complete matching of couples with no blocking pairs.

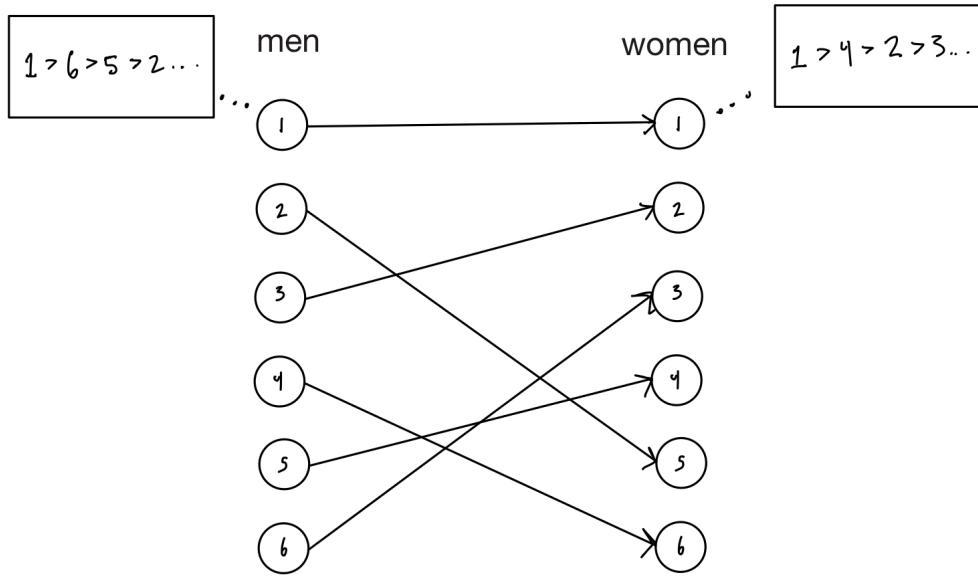


Figure 1.4: Bipartite directional graph illustrating the Stable Marriage problem

In 1962, David Gale and Lloyd Shapley published a seminal paper for the field in which they proposed the **deferred acceptance** (DA) algorithm, a solution that guarantees a stable matching for the Stable Marriage problem in polynomial time [10]. Gale and Shapley also proved their DA algorithm could solve the College Admissions problem: matching high school student applicants to colleges. College Admissions differs slightly from Stable Marriage: it is a *many-to-one* matching in which colleges accept multiple students rather than each college matching with a single student, and it has *incomplete preferences* because students do not rank (apply to) every possible college on the list. With a slight adaptation of the DA algorithm, it can solve this version of College Admissions. Since 1962, the DA algorithm has been adapted to fit many more problems with additional factors, including the residency problem.

However, there remain subsets of problems that do not fall under the same paradigm as the DA algorithm: for instance, problems with a one-sided market. The Stable Roommates problem from above is an example of a one-sided market, in which participants are matched within one undivided pool of participants rather than between two distinct categories. Stability in a one-sided market retains the same meaning as in a bipartite setting. Over 20 years after Gale and Shapley's paper, Robert Irving published a paper with a solution to the Stables Roommates problem, in which he proved his algorithm outputs a stable matching of roommates in polynomial time if one exists [6]. Irving's algorithm, similarly to the DA algorithm in bipartite markets, provides the groundwork for the majority of research on one-sided matching problems.

### 1.3.2 Fair Division

The difference between bipartite matching and fair division (also known as fair allocation) is most often *one-sided* preferences. “One-sided” preferences has a different meaning than a “one-sided” market: one-sided preferences still means a bipartite market, but only one side cares what they are matched to. Because only one side cares, we no longer have the idea that “cheating” could occur; instead, we evaluate fair division under *envy*. Say we have a cake that is half chocolate and half vanilla, and we are dividing the cake between people at a party. A person who prefers vanilla but receives chocolate will be *envious* of those who receive a vanilla slice. In a fair division problem, the goal is to minimize the amount of envy between participants and maximize the amount of *welfare*, where welfare is the benefit participants yield from the goods they receive (such as feeling significantly more joy from vanilla over chocolate).

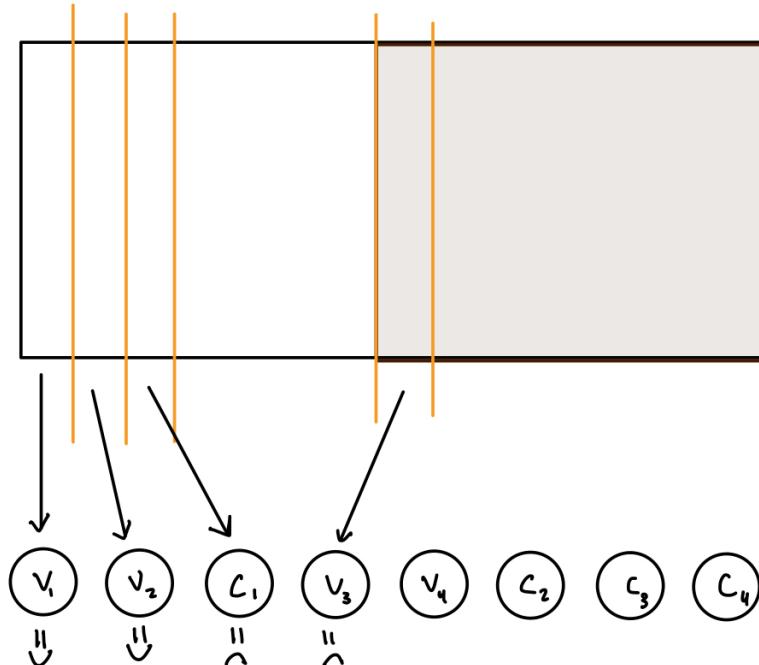


Figure 1.5: cake cutting example where  $c_1$  and  $v_3$  would envy each other

The cake-cutting example is the canon example for divisible goods [11], but there is also a lot of study into the allocation of indivisible goods: items that cannot be cut up like a cake. The aforementioned food bank example, allocating donated items (such as non-perishable cans of food) to different food banks based on preferences

they provide, involves indivisible goods [7].

In addition to envy and welfare, fair division will often be evaluated under Pareto-optimality. An allocation is considered *Pareto-optimal* if there is no other allocation that is strictly better for at least one participant while retaining the same welfare for the other participants [12]. So, we say allocation  $A$  *Pareto-dominates* allocation  $B$  if allocation  $A$  is better or the same as allocation  $B$  for every participant. As more factors are added to fair allocation problems—such as considering the order in which goods arrive for distribution, scaling the number of participants, and more specific definitions of welfare—envy-freeness and Pareto-optimality become more difficult to achieve.

## 1.4 Adding Constraints

What's known as the “vanilla” versions of these matching and division environments, which is to say the more simple or basic versions, have been studied thoroughly by now. So, to advance mechanism design research and to more accurately model real-world scenarios, scientists add *constraints* to the problems. Constraints come in many forms, and often make problems so complex they become NP-complete. In the scope of class assignment and related matching problems, constraints often come in the forms of quotas on class capacity and categories of students.

### 1.4.1 Quotas and Diversity

Quotas in allocation and matching problems can be implemented as upper and lower bounds on the mechanism. Upper bound quotas imply capacity constraints, such as a maximum number of open positions per hospital in the residency matching problem. Gale and Shapley's College Admissions problem, in which a college with  $n$  applicants can admit up to  $q$  students, is traditionally studied with these capacity constraints (where  $q$  is for quota). Lower bound quotas in something like the College Admissions would encode in the mechanism that schools need to accept at least  $m$  students of some type [13]. Strict quotas are one such constraint that can cause finding a stable matching in a bipartite setting to be NP-hard[14], so quotas are often studied under “soft” constraints as well: for example, allowing a quota to be a target range rather than a specific number.

This quota-based extension of the College Admissions problem more accurately reflects reality by acting as diversity constraints. The goal of diversity constraints is usually to produce a matching/allocation in which some type of participants is spread evenly or represented accurately in the result. In the actual college admissions problem, these constraints are commonly known as Affirmative Action, in which there are upper-bounded quotas on majority students to give space for minority students. There are many papers that mathematically study that particular implementation of diversity, as well as prove that it does not always benefit the minority applicants [15]. As a result, diversity constraints can take many forms—they have been studied under

frameworks of lower-bound quotas [13], proportionality (evenly distributing types of students in the matching) [16], and group fairness (ensuring each distinct group of students is as happy with their matching as possible based on the happiness of each individual student in the group) [17]. In each of those papers implementing diversity constraints, the constraints make the problems NP-complete.

### 1.4.2 Optimization

When we work with an NP-hard or especially an NP-complete problem, we generally do not attempt to find an algorithm that gives an exact solution; the fact that the problem is NP-hard means we don't know if it is even possible to find an exact solution efficiently. However, as we have seen with quotas and diversity, many problems with relevant real-world applications are in NP. Does that mean we simply give up trying to find a solution?

Maybe we have to let go of the idea of an exact solution, but that doesn't have to stop us from trying to find a solution we deem *good enough*. "Good enough" can mean many things depending on the problem, which is the driving force behind *optimization*. Optimization problems attempt to find a near-optimal (if not optimal) solution under given constraints. Often, these will take the form of *approximation* algorithms, which will relax some of the constraints or allow a violation of the constraints up to some upper bound to guarantee a solution that is as close to the optimal solution as mathematically possible. We can clarify what we mean by "near-optimal" based on the optimal solution: if an optimal solution has some value (such as a stable matching with no blocking pairs), we want to find a solution that maximizes/minimizes how close we get to that value (such as a matching that guarantees a maximum one blocking pair). Another optimization approach is with *heuristic* algorithms, which leans even more into the "good enough" mentality by focusing on the empirical results. Rather than making mathematical guarantees about the optimality of the solution, a heuristic algorithm will be evaluated for its performance on actual data. Usually, despite the worst-case scenario being mathematically infeasible, the algorithm is able to output a solution when given real scenarios, and therefore we consider it to be a reasonable solution. The goal of both approximation and heuristic algorithms is to find an adequate solution quickly when the optimal solution cannot be reasonably found. There are other approaches to solving optimization problems, but these approaches are the most relevant to this thesis.

## 1.5 Related Work

The problem we discuss in this thesis is related to two extensively-studied problems: school choice and course allocation. School choice in mechanism design is the problem of assigning high school students to high schools such that parents are satisfied, each student has a school to attend, and no school admits more students than they can accommodate. Course allocation in mechanism design is the problem of constructing a system to assign students to classes based on which classes they prefer to take.

### 1.5.1 Two-sided matching

In 2003, Abdulkadiroğlu and Sonmez published a seminal paper on school choice, in which they view the problem through the bipartite lens of schools and students each with distinct preferences. This lens allows them to apply the Gale-Shapley DA algorithm as a solution as well as their own algorithm, which they called Top Trading Cycles (TTC) [18]. In Top Trading Cycles, if the matching has a cycle of students in which each student prefers the next student’s school, each assignment is given to the next student until the cycle is broken.

Six years later, in 2009, Abdulkadiroğlu, Pathak, and Roth amended the 2003 paper with a case study into the New York City high school allocation mechanism, scrutinizing the DA algorithm under strategy-proofness and stability. They show empirically that there is a loss of Pareto efficiency in the algorithm when adding strategy-proofness and stability to the environment [19].

Many papers, including some outside of school choice specifically, extend one or both of Abdulkadiroğlu’s papers in further research. In 2010, Biró et al. studied school choice in a college setting with upper and lower quotas, constraining the environment in a new way. They determined that stable matchings do not always exist with those quotas, and in fact determining whether a stable matching exists at all is NP-complete [13]. Three years later, Hafalir et al. analyzed the DA algorithm for school choice under quotas for diversity, where they found that using upper quotas produced a matching that was Pareto-dominated by a matching with lower quotas for minority students [15]. There has since been much more research into the stability effects of quotas on the two-sided school choice problem, including papers that implicate stability with diverse students of multiple types [20].

### 1.5.2 One-sided matching

Since 2003, research also branched from two-sided school choice into one-sided matching, where only one side of the market has preferences. This was often studied under the framework of course allocation, a mechanism design problem which arose from the existing school choice literature. In 2011, Budish and Cantillon published the seminal paper on course allocation at Harvard (which cites Abdulkadiroğlu’s 2003 and 2009 papers) in which they evaluate Harvard’s mechanism under strategyproofness and pareto-optimality [21]. There have been slight alterations to their framework since then, such as the effects of incomplete preferences on participant welfare.

Only in recent years have scholars started studying one-sided matching with quotas and other diversity constraints. In fact, the constraints closest to that in this paper (groups of students as well as sub-types within the groups) only seem to exist in literature published as recently as 2025. Specifically, Santhini et al. researched the empirical effects of both soft quotas [14] and group fairness [22] on the mechanism using convex optimization (also done similarly by Panda et al. [17]). There is not readily as much theory-based research on one-sided matching with constraints as two-sided matching, and those that do study it under a theoretical framework use more

of a fair division analysis than stability.

As far as we can tell, the closest paper to our construction of stability in this thesis is an unpublished 2011 paper by Abdulkadiroğlu, who generalizes the two-sided and one-sided settings to one environment where the TTC algorithm can be reduced to work for both kinds of preferences. In the paper, he frames stability to be “if, whenever a student prefers a school to her assignment, that school is enrolled up to its capacity by students who have higher priority at that school” [23], where “higher priority” in the one-sided setting only considers how strongly the other students prefer that school. The paper does not consider the effect of adding diversity constraints.